

Solving Kernel Ridge Regression with Gradient Descent for a Non-Constant Kernel (after an introduction to kernels)

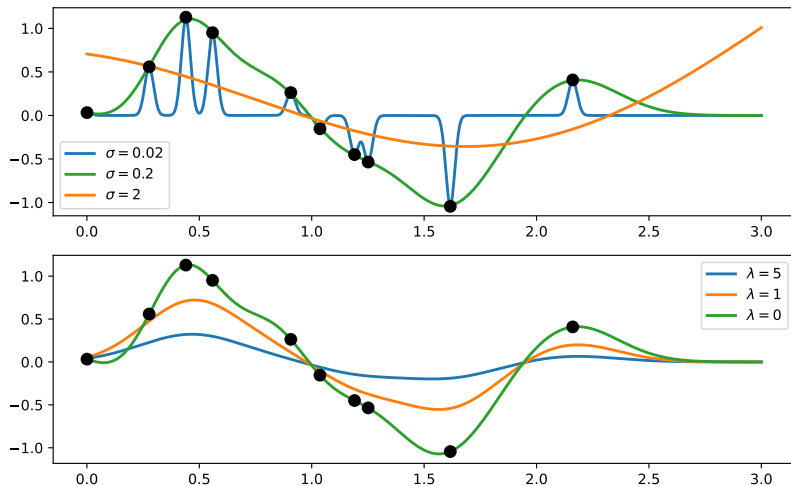
Oskar Allerbo

KTH Royal Institute of Technology

2024-03-11

www.github.com/allerbo/hemavan24/slides.pdf

Kernel Ridge Regression



Kernel Ridge Regression

Kernel ridge regression (KRR) is a generalization of linear ridge regression that is

- non-linear in the data,
- but linear in the parameters.
- a convex problem, with a closed-form solution.

Outline

- 1 Introduction to Kernels
- 2 Kernel Principal Component Analysis (KPCA)
- 3 Kernel Ridge Regression (KRR)
- 4 Kernel Gradient Descent for Non-constant Kernels (KGD)
- 5 Generalization to Neural Networks (NNs)
- 6 Conclusions

What is a Kernel?

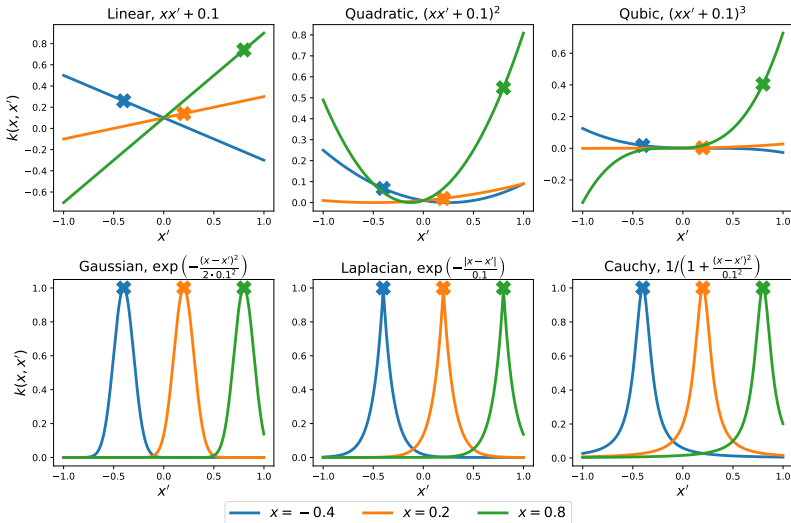
A kernel function:

- Takes two arguments and outputs a scalar:
 $k(\mathbf{x}, \mathbf{x}') \in \mathbb{R}$. $\mathbf{x}, \mathbf{x}' \in \mathbb{R}^p$.
- Is symmetric:
 $k(\mathbf{x}, \mathbf{x}') = k(\mathbf{x}', \mathbf{x})$.
- Is positive semi-definite:
 $\sum_{i=1}^n \sum_{j=1}^n c_i c_j k(\mathbf{x}_i, \mathbf{x}_j) \geq 0$. For all $\mathbf{x}_i, \mathbf{x}_j \in \mathbb{R}^p$, $c_i, c_j \in \mathbb{R}$.
- Is the dot product of the feature expansions of \mathbf{x} and \mathbf{x}' :
 $k(\mathbf{x}, \mathbf{x}') = \varphi(\mathbf{x})^\top \varphi(\mathbf{x}')$. $\varphi(\mathbf{x}), \varphi(\mathbf{x}') \in \mathbb{R}^q$.
 $q = \infty$ is a possibility!
- (Is associated to a Reproducing Kernel Hilbert Space, RKHS.)

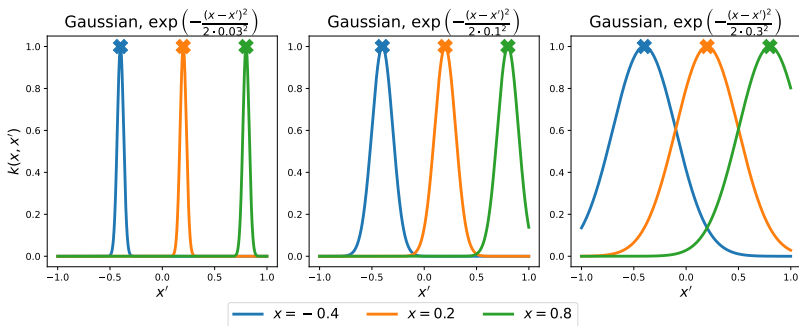
Examples of Kernels

- Linear: $k(\mathbf{x}, \mathbf{x}') = \mathbf{x}^\top \mathbf{x}' + c, c \in \mathbb{R}$
- Polynomial: $k(\mathbf{x}, \mathbf{x}') = (\mathbf{x}^\top \mathbf{x}' + c)^q, c \in \mathbb{R}, q \in \mathbb{N}$
- Gaussian: $k(\mathbf{x}, \mathbf{x}') = e^{-\frac{\|\mathbf{x} - \mathbf{x}'\|_2^2}{2\sigma^2}}, \sigma > 0$
- Laplace: $k(\mathbf{x}, \mathbf{x}') = e^{-\frac{\|\mathbf{x} - \mathbf{x}'\|_2}{\sigma}}, \sigma > 0$
- Cauchy: $k(\mathbf{x}, \mathbf{x}') = \frac{1}{1 + \frac{\|\mathbf{x} - \mathbf{x}'\|_2^2}{\sigma^2}}, \sigma > 0$
- Matérn: $k(\mathbf{x}, \mathbf{x}') = \frac{2^{1-\nu}}{\Gamma(\nu)} \cdot \left(\sqrt{2\nu} \cdot \frac{\|\mathbf{x} - \mathbf{x}'\|_2}{\sigma} \right)^\nu \cdot K_\nu \left(\sqrt{2\nu} \cdot \frac{\|\mathbf{x} - \mathbf{x}'\|_2}{\sigma} \right)$

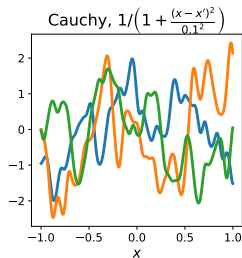
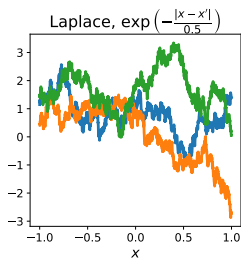
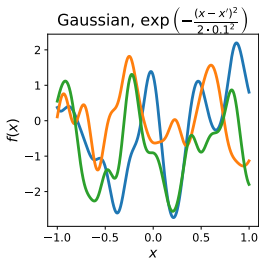
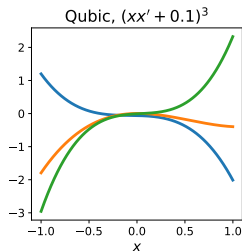
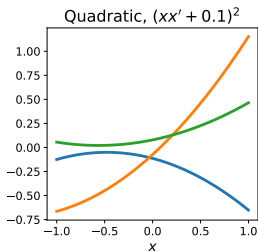
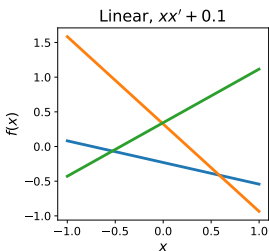
Examples of Kernels (1D Case)



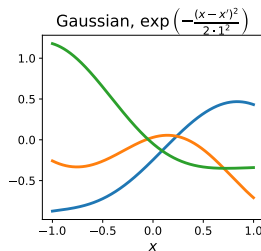
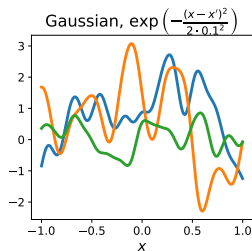
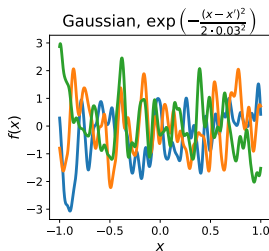
Examples of Kernels (1D Case)



Examples of Functions (1D Case)



Examples of Functions (1D Case)



Examples of Feature Expansions (1D Case)

- Linear, $k(x, x') = x \cdot x' + c$

$$\varphi(x) = [x \quad \sqrt{c}]^\top$$

$$\varphi(x)^\top \varphi(x') = x \cdot x' + \sqrt{c} \cdot \sqrt{c}$$

- Quadratic, $k(x, x') = (x \cdot x' + c)^2$

$$\varphi(x) = [x^2 \quad \sqrt{2c}x \quad c]^\top$$

$$\varphi(x)^\top \varphi(x') = x^2 \cdot x'^2 + 2c \cdot x \cdot x' + c \cdot c$$

- Gaussian, $k(x, x') = e^{-\frac{(x-x')^2}{2\sigma^2}}$

$$\varphi(x) = e^{-\frac{x^2}{2\sigma^2}} \left[1 \quad \frac{x^1}{\sigma^1 \sqrt{1!}} \quad \dots \quad \frac{x^k}{\sigma^k \sqrt{k!}} \dots \right]^\top$$

$$\varphi(x)^\top \varphi(x') = e^{-\frac{1}{2\sigma^2}(x^2 + x'^2)} \cdot \underbrace{\sum_{k=0}^{\infty} \frac{(x \cdot x' / \sigma^2)^k}{k!}}_{= e^{\frac{2 \cdot x \cdot x'}{2\sigma^2}}}$$

Notation

Training Data: $\mathbf{X} \in \mathbb{R}^{n \times p}$, $\mathbf{y} \in \mathbb{R}^n$.

New Observations: $\mathbf{X}^* \in \mathbb{R}^{n^* \times p}$.

Predictions: $\hat{\mathbf{f}} \in \mathbb{R}^n$, $\hat{\mathbf{f}}^* \in \mathbb{R}^{n^*}$.

Feature Expansion Matrices:

$$\Phi = \Phi(\mathbf{X}) \in \mathbb{R}^{n \times q}, \Phi^* = \Phi^*(\mathbf{X}^*) \in \mathbb{R}^{n^* \times q}.$$

Kernel Matrices:

$$\mathbf{K} = \mathbf{K}(\mathbf{X}) \in \mathbb{R}^{n \times n}, \mathbf{K}^* = \mathbf{K}(\mathbf{X}^*, \mathbf{X}) \in \mathbb{R}^{n^* \times n}.$$

$$\mathbf{K} = \Phi \Phi^\top, \mathbf{K}^* = \Phi^* \Phi^\top.$$

$$k(\mathbf{x}, \mathbf{x}') = \varphi(\mathbf{x})^\top \varphi(\mathbf{x}').$$

Notation

$$X = \begin{bmatrix} - & x_1^\top & - \\ - & x_2^\top & - \\ \vdots & \vdots & \vdots \\ - & x_n^\top & - \end{bmatrix}$$

$$X^* = \begin{bmatrix} - & x_1^{*\top} & - \\ - & x_2^{*\top} & - \\ \vdots & \vdots & \vdots \\ - & x_{n^*}^{*\top} & - \end{bmatrix}$$

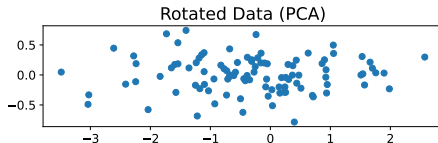
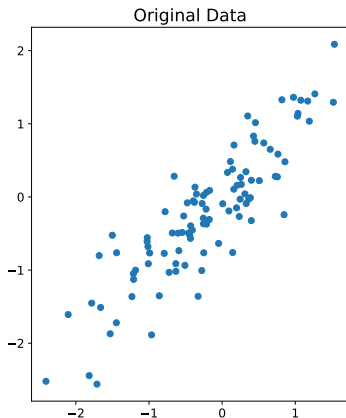
$$\Phi = \begin{bmatrix} - & \varphi(x_1)^\top & - \\ - & \varphi(x_2)^\top & - \\ \vdots & \vdots & \vdots \\ - & \varphi(x_n)^\top & - \end{bmatrix}$$

$$\Phi^* = \begin{bmatrix} - & \varphi(x_1^*)^\top & - \\ - & \varphi(x_2^*)^\top & - \\ \vdots & \vdots & \vdots \\ - & \varphi(x_{n^*}^*)^\top & - \end{bmatrix}$$

$$\Phi^* \Phi^\top = K^* = \begin{bmatrix} k(x_1^*, x_1) & k(x_1^*, x_2) & \dots & k(x_1^*, x_n) \\ k(x_2^*, x_1) & k(x_2^*, x_2) & \dots & k(x_2^*, x_n) \\ \vdots & \vdots & \ddots & \vdots \\ k(x_{n^*}^*, x_1) & k(x_{n^*}^*, x_2) & \dots & k(x_{n^*}^*, x_n) \end{bmatrix}$$

Kernel Principal Component Analysis

Principal Component Analysis (PCA):
Rotate data to find directions with maximum variance.



Kernel Principal Component Analysis

Principal Component Analysis (PCA):

Rotate data to find directions with maximum variance.

$$\mathbf{X}^\top \mathbf{X} = \mathbf{P} \mathbf{D} \mathbf{P}^\top, \mathbf{Z} = \mathbf{X} \mathbf{P}.$$

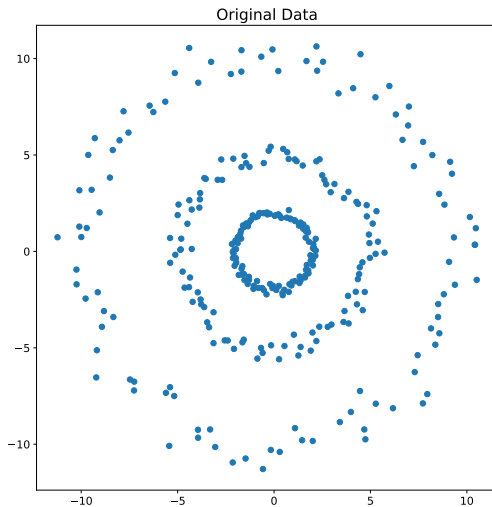
Equivalent, dual formulation:

$$\mathbf{X} \mathbf{X}^\top = \mathbf{U} \mathbf{D} \mathbf{U}^\top, \mathbf{Z} = \mathbf{U} \sqrt{\mathbf{D}}.$$

Kernel PCA:

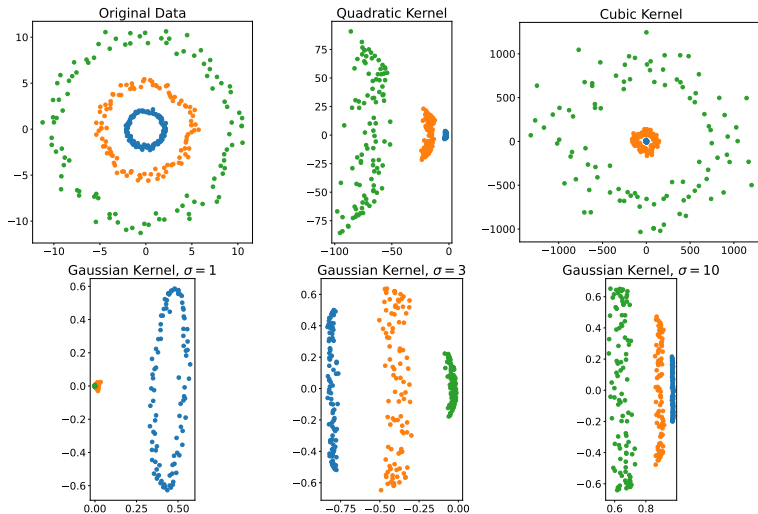
$$\Phi \Phi^\top = \mathbf{K} = \mathbf{U}_K \mathbf{D}_K \mathbf{U}_K^\top, \mathbf{Z}_K = \mathbf{U}_K \sqrt{\mathbf{D}_K}.$$

Kernel Principal Component Analysis

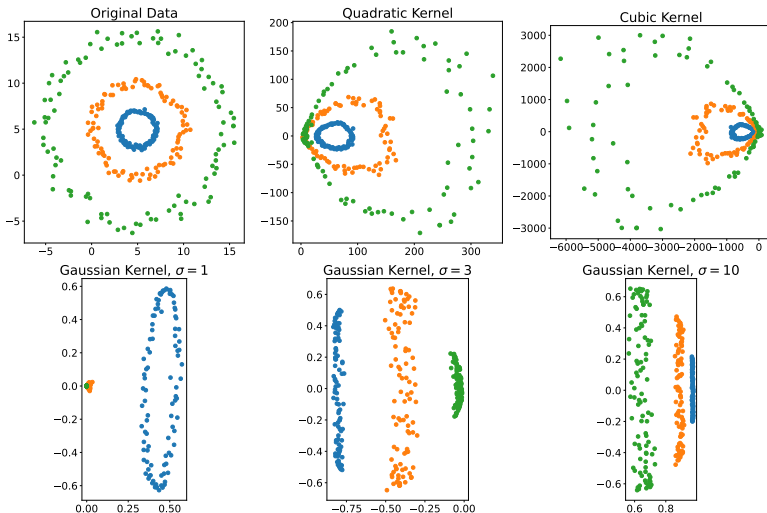


What about this?

Kernel Principal Component Analysis



Kernel Principal Component Analysis



Kernel Ridge Regression

- Linear Ridge Regression
- Dual Formulation of Linear Ridge Regression
- Ridge Regression in Feature Space
- Dual Formulation of Ridge Regression in Feature Space
(=Kernel Ridge Regression)

Linear Ridge Regression

$$\begin{aligned}\hat{\beta} &= \operatorname{argmin}_{\beta \in \mathbb{R}^p} \frac{1}{2} \|\mathbf{y} - \mathbf{X}\beta\|_2^2 + \frac{\lambda}{2} \underbrace{\|\beta\|_2^2}_{=\beta^\top \beta} \\ &= (\mathbf{X}^\top \mathbf{X} + \lambda \mathbf{I}_{p \times p})^{-1} \mathbf{X}^\top \mathbf{y}\end{aligned}$$

Predictions are given by

$$\begin{bmatrix} \hat{\mathbf{f}} \\ \hat{\mathbf{f}}^* \end{bmatrix} = \begin{bmatrix} \mathbf{X} \hat{\beta} \\ \mathbf{X}^* \hat{\beta} \end{bmatrix} = \begin{bmatrix} \mathbf{X} (\mathbf{X}^\top \mathbf{X} + \lambda \mathbf{I}_{p \times p})^{-1} \mathbf{X}^\top \mathbf{y} \\ \mathbf{X}^* (\mathbf{X}^\top \mathbf{X} + \lambda \mathbf{I}_{p \times p})^{-1} \mathbf{X}^\top \mathbf{y} \end{bmatrix}.$$

Dual Formulation of Linear Ridge Regression

Dual formulation for $\beta = \mathbf{X}^\top \alpha$:

$$\begin{aligned}\hat{\alpha} &= \underset{\alpha \in \mathbb{R}^n}{\operatorname{argmin}} \frac{1}{2} \|\mathbf{y} - \mathbf{X}\mathbf{X}^\top \alpha\|_2^2 + \frac{\lambda}{2} \underbrace{\alpha^\top \mathbf{X}\mathbf{X}^\top \alpha}_{=\|\alpha\|_{\mathbf{X}\mathbf{X}^\top}^2} \\ &= (\mathbf{X}\mathbf{X}^\top + \lambda \mathbf{I}_{n \times n})^{-1} \mathbf{y}\end{aligned}$$

Predictions are given by

$$\begin{bmatrix} \hat{\mathbf{f}} \\ \hat{\mathbf{f}}^* \end{bmatrix} = \begin{bmatrix} \mathbf{X} \cdot \mathbf{X}^\top \hat{\alpha} \\ \mathbf{X}^* \cdot \mathbf{X}^\top \hat{\alpha} \end{bmatrix} = \begin{bmatrix} \mathbf{X}\mathbf{X}^\top (\mathbf{X}\mathbf{X}^\top + \lambda \mathbf{I}_{n \times n})^{-1} \mathbf{y} \\ \mathbf{X}^* \mathbf{X}^\top (\mathbf{X}\mathbf{X}^\top + \lambda \mathbf{I}_{n \times n})^{-1} \mathbf{y} \end{bmatrix}.$$

Linear Ridge Regression

Predictions given by

$$\begin{bmatrix} \mathbf{X}\hat{\boldsymbol{\beta}} \\ \mathbf{X}^*\hat{\boldsymbol{\beta}} \end{bmatrix} = \begin{bmatrix} \mathbf{X}(\mathbf{X}^\top \mathbf{X} + \lambda \mathbf{I}_{p \times p})^{-1} \mathbf{X}^\top \mathbf{y} \\ \mathbf{X}^*(\mathbf{X}^\top \mathbf{X} + \lambda \mathbf{I}_{p \times p})^{-1} \mathbf{X}^\top \mathbf{y} \end{bmatrix}$$

and

$$\begin{bmatrix} \mathbf{X}\mathbf{X}^\top \hat{\boldsymbol{\alpha}} \\ \mathbf{X}^*\mathbf{X}^\top \hat{\boldsymbol{\alpha}} \end{bmatrix} = \begin{bmatrix} \mathbf{X}\mathbf{X}^\top (\mathbf{X}\mathbf{X}^\top + \lambda \mathbf{I}_{n \times n})^{-1} \mathbf{y} \\ \mathbf{X}^*\mathbf{X}^\top (\mathbf{X}\mathbf{X}^\top + \lambda \mathbf{I}_{n \times n})^{-1} \mathbf{y} \end{bmatrix}.$$

However,

$$(\mathbf{X}^\top \mathbf{X} + \lambda \mathbf{I}_{p \times p})^{-1} \mathbf{X}^\top = \mathbf{X}^\top (\mathbf{X}\mathbf{X}^\top + \lambda \mathbf{I}_{n \times n})^{-1}.$$

Ridge Regression in Feature Space

$$\mathbf{x} \in \mathbb{R}^p \mapsto \varphi(\mathbf{x}) \in \mathbb{R}^q.$$

E.g. polynomial regression, $x \mapsto [1, x, x^2, \dots, x^{q-1}]$.

$$\mathbf{X} \in \mathbb{R}^{n \times p} \mapsto \Phi \in \mathbb{R}^{n \times q}, \mathbf{X}^* \in \mathbb{R}^{n^* \times p} \mapsto \Phi^* \in \mathbb{R}^{n^* \times q}$$

$$\hat{\beta} = \operatorname{argmin}_{\beta \in \mathbb{R}^q} \frac{1}{2} \|\mathbf{y} - \Phi\beta\|_2^2 + \frac{\lambda}{2} \|\beta\|_2^2$$

$$\hat{\alpha} = \operatorname{argmin}_{\alpha \in \mathbb{R}^n} \frac{1}{2} \|\mathbf{y} - \Phi\Phi^\top \alpha\|_2^2 + \frac{\lambda}{2} \|\alpha\|_{\Phi\Phi^\top}^2$$

Predictions are given by

$$\begin{bmatrix} \hat{\mathbf{f}} \\ \hat{\mathbf{f}}^* \end{bmatrix} = \begin{bmatrix} \Phi (\Phi^\top \Phi + \lambda I_{p \times p})^{-1} \Phi^\top \mathbf{y} \\ \Phi^* (\Phi^\top \Phi + \lambda I_{p \times p})^{-1} \Phi^\top \mathbf{y} \end{bmatrix} = \begin{bmatrix} \Phi\Phi^\top (\Phi\Phi^\top + \lambda I_{n \times n})^{-1} \mathbf{y} \\ \Phi^*\Phi^\top (\Phi\Phi^\top + \lambda I_{n \times n})^{-1} \mathbf{y} \end{bmatrix}$$

Kernel Ridge Regression

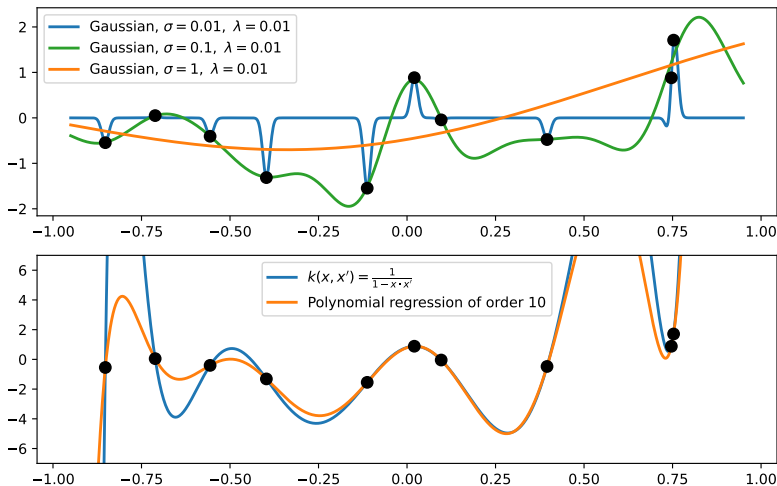
For $\mathbf{K} = \Phi\Phi^\top \in \mathbb{R}^{n \times n}$, $\mathbf{K}^* = \Phi^*\Phi^\top \in \mathbb{R}^{n^* \times n}$,

$$\hat{\alpha} = \operatorname{argmin}_{\alpha \in \mathbb{R}^n} \frac{1}{2} \|\mathbf{y} - \mathbf{K}\alpha\|_2^2 + \frac{\lambda}{2} \|\alpha\|_{\mathbf{K}}^2$$

$$\begin{bmatrix} \hat{f} \\ \hat{f}^* \end{bmatrix} = \begin{bmatrix} \mathbf{K} \\ \mathbf{K}^* \end{bmatrix} \hat{\alpha} = \begin{bmatrix} \mathbf{K}(\mathbf{K} + \lambda \mathbf{I}_{n \times n})^{-1} \mathbf{y} \\ \mathbf{K}^*(\mathbf{K} + \lambda \mathbf{I}_{n \times n})^{-1} \mathbf{y} \end{bmatrix}$$

$q = \infty$ is OK, since Φ and Φ^* are never explicitly calculated.

Kernel Ridge Regression



$$\frac{1}{1-x} = 1 + x + x^2 + x^3 + x^4 + \dots, \quad |x| < 1$$

Kernel Gradient Descent for Non-Constant Kernels

Kernel gradient descent in function/prediction space:
(with regularization through early stopping)

$$\begin{bmatrix} \hat{\mathbf{f}}_{t+1} \\ \hat{\mathbf{f}}_{t+1}^* \end{bmatrix} = \begin{bmatrix} \hat{\mathbf{f}}_t \\ \hat{\mathbf{f}}_t^* \end{bmatrix} - \eta \cdot \begin{bmatrix} \mathbf{K} \\ \mathbf{K}^* \end{bmatrix} (\hat{\mathbf{f}}_t - \mathbf{y})$$

$$\left(\text{where } \begin{bmatrix} \hat{\mathbf{f}} \\ \hat{\mathbf{f}}^* \end{bmatrix} = \begin{bmatrix} \mathbf{K} \\ \mathbf{K}^* \end{bmatrix} \hat{\alpha} \right)$$

Time dependent kernels:

$$\begin{bmatrix} \hat{\mathbf{f}}_{t+1} \\ \hat{\mathbf{f}}_{t+1}^* \end{bmatrix} = \begin{bmatrix} \hat{\mathbf{f}}_t \\ \hat{\mathbf{f}}_t^* \end{bmatrix} - \eta \cdot \begin{bmatrix} \mathbf{K}_t \\ \mathbf{K}_t^* \end{bmatrix} (\hat{\mathbf{f}}_t - \mathbf{y})$$

Kernel Regression with Gradient Descent and Non-Constant Kernels

Proposition

For a translational invariant kernel with bandwidth σ ,
 $k(\mathbf{x}, \mathbf{x}', \sigma) = k\left(\frac{\|\mathbf{x} - \mathbf{x}'\|_2}{\sigma}\right)$, and for constant training time, t ,

$$\left\| \nabla_{\mathbf{x}^*} \hat{\mathbf{f}}(\mathbf{x}^*, t) \right\|_2 \leq \frac{1}{\sigma} \cdot t \cdot C(k(\cdot), \mathbf{X}, \mathbf{y})$$

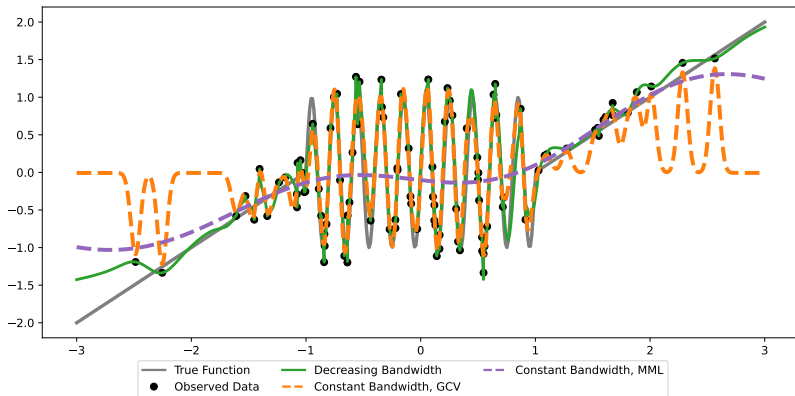
i.e. a larger bandwidth results in a simpler model, and a smaller bandwidth in a more complex model.

We use $1/\sigma$ as a proxy for complexity.

$$\begin{bmatrix} \hat{\mathbf{f}}_{t+1} \\ \hat{\mathbf{f}}_{t+1}^* \end{bmatrix} = \begin{bmatrix} \hat{\mathbf{f}}_t \\ \hat{\mathbf{f}}_t^* \end{bmatrix} - \eta \cdot \begin{bmatrix} \mathbf{K}(\sigma_t) \\ \mathbf{K}^*(\sigma_t) \end{bmatrix} (\hat{\mathbf{f}}_t - \mathbf{y})$$

Idea: Start with a kernel with large bandwidth (a simple model).
Gradually decrease the bandwidth towards zero during training.

Kernel Regression with Gradient Descent and Non-Constant Kernels

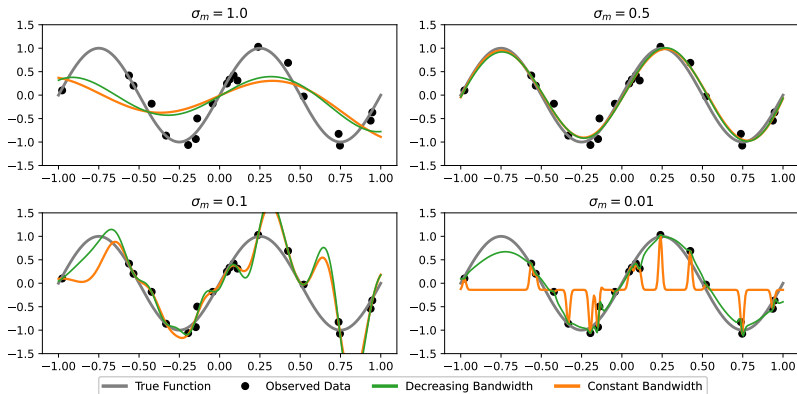


Kernel Regression with Gradient Descent and Non-Constant Kernels, Double Descent

Generalization as function of model complexity:

A too simple model generalizes poorly.	}	Classical statistical knowledge
A model of appropriate complexity generalizes well.		
A too complex model generalizes poorly (overfitting).		
An extremely complex model generalizes well (double descent).		

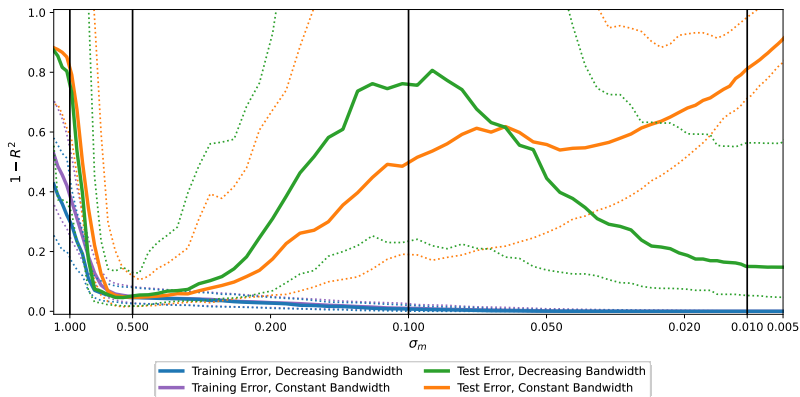
Kernel Regression with Gradient Descent and Non-Constant Kernels, Double Descent



σ_m : Minimum allowed bandwidth when decreasing the bandwidth.

Employed bandwidth when using a constant bandwidth.

Kernel Regression with Gradient Descent and Non-Constant Kernels, Double Descent



Kernel Regression with Gradient Descent and Non-Constant Kernels, Double Descent

Proposition (Simplified)

$$|\hat{f}(\mathbf{x}^*, t, \sigma_m)| \leq \min \left(\overline{\sigma^{-1}}(\sigma_m), \overline{k_2^*}(\sigma_m) \right) \cdot t,$$

where $\overline{\sigma^{-1}}(\sigma_m)$ increases with model complexity and $\overline{k_2^*}(\sigma_m)$ decreases with model complexity.

Low complexity:	$\overline{\sigma^{-1}}(\sigma_m)$ is small	$ \hat{f} $ is too small.
Moderate complexity:	$\overline{\sigma^{-1}}(\sigma_m)$ is moderate	$ \hat{f} $ is appropriate.
High complexity:	$\overline{\sigma^{-1}}(\sigma_m)$ is large	$ \hat{f} $ is too large.
Very high complexity:	$\overline{k_2^*}(\sigma_m)$ is moderate	$ \hat{f} $ is appropriate.

Neural Tangent Kernel Gradient Flow

$$\begin{aligned} \begin{bmatrix} \hat{\mathbf{f}}(t + \Delta t) \\ \hat{\mathbf{f}}^*(t + \Delta t) \end{bmatrix} &= \begin{bmatrix} \hat{\mathbf{f}}(t) \\ \hat{\mathbf{f}}^*(t) \end{bmatrix} - \Delta t \cdot \begin{bmatrix} \mathbf{K}(t) \\ \mathbf{K}^*(t) \end{bmatrix} (\hat{\mathbf{f}}(t) - \mathbf{y}) \\ \Leftrightarrow \frac{\begin{bmatrix} \hat{\mathbf{f}}(t + \Delta t) \\ \hat{\mathbf{f}}^*(t + \Delta t) \end{bmatrix} - \begin{bmatrix} \hat{\mathbf{f}}(t) \\ \hat{\mathbf{f}}^*(t) \end{bmatrix}}{\Delta t} &= - \begin{bmatrix} \mathbf{K}(t) \\ \mathbf{K}^*(t) \end{bmatrix} (\hat{\mathbf{f}}(t) - \mathbf{y}) \end{aligned}$$

$$\begin{aligned} \begin{bmatrix} \frac{\partial \hat{\mathbf{f}}(t)}{\partial t} \\ \frac{\partial \hat{\mathbf{f}}^*(t)}{\partial t} \end{bmatrix} &\stackrel{\text{C.R.}}{=} \begin{bmatrix} \frac{\partial \hat{\mathbf{f}}(t)}{\partial \hat{\boldsymbol{\theta}}(t)} \\ \frac{\partial \hat{\mathbf{f}}^*(t)}{\partial \hat{\boldsymbol{\theta}}(t)} \end{bmatrix} \cdot \frac{\partial \hat{\boldsymbol{\theta}}(t)}{\partial t} \stackrel{\text{G.D.}}{=} - \begin{bmatrix} \frac{\partial \hat{\mathbf{f}}(t)}{\partial \hat{\boldsymbol{\theta}}(t)} \\ \frac{\partial \hat{\mathbf{f}}^*(t)}{\partial \hat{\boldsymbol{\theta}}(t)} \end{bmatrix} \cdot \frac{\partial L(\hat{\mathbf{f}}(t))}{\partial \hat{\boldsymbol{\theta}}(t)} \\ &\stackrel{\text{C.R.}}{=} - \begin{bmatrix} \frac{\partial \hat{\mathbf{f}}(t)}{\partial \hat{\boldsymbol{\theta}}(t)} \\ \frac{\partial \hat{\mathbf{f}}^*(t)}{\partial \hat{\boldsymbol{\theta}}(t)} \end{bmatrix} \cdot \left(\frac{\partial \hat{\mathbf{f}}(t)}{\partial \hat{\boldsymbol{\theta}}(t)} \right)^\top \cdot \frac{\partial L(\hat{\mathbf{f}}(t))}{\partial \hat{\mathbf{f}}(t)} \\ &\stackrel{\text{def}}{=} - \begin{bmatrix} \boldsymbol{\Phi}(t) \boldsymbol{\Phi}(t)^\top \\ \boldsymbol{\Phi}^*(t) \boldsymbol{\Phi}(t)^\top \end{bmatrix} \cdot \frac{\partial L(\hat{\mathbf{f}}(t))}{\partial \hat{\mathbf{f}}(t)} = - \begin{bmatrix} \mathbf{K}(t) \\ \mathbf{K}^*(t) \end{bmatrix} \cdot \frac{\partial L(\hat{\mathbf{f}}(t))}{\partial \hat{\mathbf{f}}(t)}. \end{aligned}$$

Improved Generalization by Neural Tangent Control

- Kernel gradient descent performs best when the bandwidth decreases toward zero,
- that is, when $\begin{bmatrix} K \\ K^* \end{bmatrix}$ goes to $\begin{bmatrix} I_{n \times n} \\ \mathbf{0}_{n^* \times n} \end{bmatrix}$.
- Can this behaviour be enforced for the neural tangent kernel?
- Work in progress, but it seems so...

Conclusions

Kernels models

- are non-linear in data.
- are linear in parameters (and thus convex with closed-form solutions).
- correspond to a (possibly infinite dimensional) feature expansion.

Examples:

- Kernel Principal Component Analysis.
- Kernel Ridge Regression.

Kernel regression with gradually increased model complexity

- reduces the need for hyper parameter selection.
- exhibits a double descent behavior.
- is generalizable to any parametric model trained with gradient descent.

The End

Thank you!